Names: Mike Pauls, Jayati Singh, Austin Lee
NetIds: mepauls2, jayati, sal3
Team name: team_name
School: On-campus

Kernels (>90% of program time):

32.06% 35.969ms 20 1.7984ms 1.1200us 33.609ms [CUDA memcpy HtoD]
17.88% 20.062ms 1 20.062ms 20.062ms 20.062ms volta_scudnn_128x64_relu_interior_nn_v1
17.16% 19.252ms 4 4.8129ms 4.8122ms 4.8133ms volta_gcgemm_64x32_nt
8.53% 9.5671ms 4 2.3918ms 2.0052ms 3.1255ms void fft2d_c2r_32x32<float, bool=0, bool=0, unsigned int=0, bool=0, bool=0>(float*, float2 const *, int, int, int, int, int, int, int, int, float, float, cudnn::reduced_divisor, bool, float*, float*, int2, int, int)
7.80% 8.7556ms 1 8.7556ms 8.7556ms 8.7556ms volta_sgemm_128x128_tn
6.42% 7.2052ms 2 3.6026ms 25.536us 7.1797ms void op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7, cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*, cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float, dimArray, reducedDivisorArray)
5.70% 6.3895ms 4 1.5974ms 1.2742ms 2.0207ms void fft2d_r2c_32x32<float, bool=0, unsigned int=0, bool=0>(float2*, float const *, int, int, int, int, int, int, int, int, cudnn::reduced_divisor, bool, int2, int, int)
3.88% 4.3527ms 1 4.3527ms 4.3527ms 4.3527ms void cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)

CUDA API Calls (>90% program time):

41.41% 3.08766s 22 140.35ms 14.396us 1.61488s cudaStreamCreateWithFlags
33.15% 2.47141s 24 102.98ms 55.402us 2.46633s cudaMemGetInfo
21.17% 1.57836s 19 83.072ms 1.2440us 421.47ms cudaFree

Kernels vs. API:

CUDA API calls are instructions (cudaMemcpy, cudaGetDevice, etc.) that are used by the host to create and communicate with the kernels, while GPU Activities (kernels) are the actual instructions in the kernel that run during each invocation.

Rai running MXNet on the CPU:

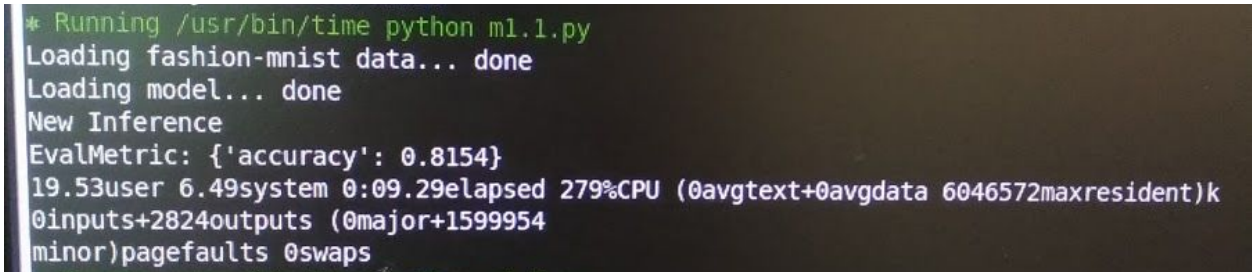

Program run time:
User: 19.53 seconds
System: 6.49 seconds
Elapsed: 0:09.29

Rai running MXNet on the GPU:

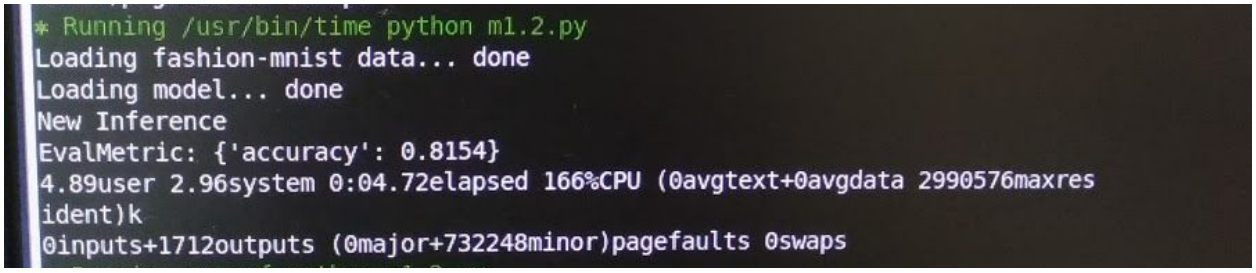

Program run time:
User: 4.89 seconds
System: 2.96 seconds
Elapsed: 0:04.72

Whole program execution time:
<u>New Inference:10000</u>
User: 88.36 seconds
System: 10.38 seconds
Elapsed: 1:16.79 seconds
Op Time: 11.134082 seconds
Op Time: 61.390580 seconds
Correctness: 0.7653 Model: ece408

<u>New Inference:1000</u>
User: 18.35 seconds
System: 2.70 seconds
Elapsed: 0:11.22
Op Time: 1.317549
Op Time: 6.760934
Correctness: 0.767 Model: ece408

New Inference:100
User: 8.61
System: 2.60
Elapsed: 0:03.16
Op Time: 0.119225
Op Time: 0.676391
Correctness: 0.76 Model: ece408